




Machine learning in a nutshell

Machine learning uses algorithms to identify patterns within data. Patterns that are used to create a **data model that can make predictions**. The **Azure for the Data Scientist course (DP-100)** focuses on creating and using machine learning models with **Azure Machine Learning**. To understand the purpose of the exercises, a simplified overview is provided here.

1.


Define the problem

Decide on what the model should predict and when it is successful.




Classification

Predict a categorical value.



Regression

Predict a numerical value.




Time-series forecasting

Predict future numerical values based on time-series data.



Computer Vision

Classify images or detect objects in images.




Natural Language Processing

Extract insights from text.

2.


Get the data

Find data sources and get access. Azure Machine Learning connects seamlessly with the three most used Azure storages for data science.




Azure Blob Storage

Object cloud storage. Uses flat namespace to store unstructured data.



Azure Data Lake Gen2

Unlimited object cloud storage. Uses hierarchical namespace for granular access control.



Azure SQL Database

Relational cloud database. Used for tabular and transactional data.

3.


Prepare the data

Explore the data. Clean and transform the data based on the model's requirements.




Exploratory data analysis (EDA)

Analyse your data, get summary statistics, and understand possible correlations between variables.



Feature engineering:

Transform the data to create features that will help the model to predict the target value.



Create validation set: Split the data into training and validation or test dataset to evaluate the model.

4.


Train the model

Choose algorithm and hyperparameters based on trial and error.




Data

Includes features (what influences the to be predicted value) and the target value (if it exists).



Algorithm

Based on the task (e.g. classification), different [algorithms](#) and hyperparameters can be tried.




Model

Often stored as a binary file (e.g. pickle file). Use on new data with same features to predict the target value.

5.


Integrate the model

Use endpoint to generate predictions.



Real-time predictions

Create light-weight app to predict target value in real-time per new data measurement.



Batch predictions

Create pipeline to predict target value on new set of data measurements.

6.

Monitor the model

Track the model's performance.



Data drift

When new data differs significantly from training dataset.



Evaluation metrics

Keep track of the model's performance. When the model's predictions are increasingly incorrect.

Retrain the model

Learn more on:

[Create machine learning models](#)

[Explore visual tools for machine learning on Microsoft Learn](#)

[Self-paced Microsoft Learn content on using Azure Machine Learning](#)

[Azure Machine Learning documentation](#)